

Automatic Learning of Summary Statistics for Approximate Bayesian Computation Using Deep Learning

Samuel Wqvist, Ph.D. Student, Lund University.

Work with: Pierre-Alexandre Mattei (ITU), Umberto Picchini (GU/Chalmers), and Jes Frellsen (ITU).

What we will talk about today

- Introduction to ABC;
- How to leverage deep learning methods to learn the summary statistics for ABC;
 - Present the main results from the paper: *Partially Exchangeable Networks and Architectures for Learning Summary Statistics in Approximate Bayesian Computation* (accepted for ICML 2019);
- We will have a practical focus and run ABC for a simple model (the Beta-Binomial model).

Approximate Bayesian Computation: Simulation based inference

- ABC in a nut-shell: Simulations-based inference method where we generate parameter proposals θ^* and accept θ^* if the generated data $y^* \sim p(y|\theta^*)$ is *similar* to our observed data y^{obs} ;
- ABC only requires that we can simulate data from a computer simulator of our model $p(y|\theta)$.
- Thus ABC is very generic, and can be applied for models where the likelihood function is intractable.

- *Curse-of-dimensionality*: Instead of comparing the data sets we compare a set of summary statistics $s = S(y)$. The main focus of our work is how to *automatically* learn the summary statistics. (For example for dynamic models, summaries can be autocorrelations, cross-covariances, stationary mean. For i.i.d. data could be quantiles, mean and standard deviation etc.;)

Approximate Bayesian Computation: Rejection sampling method

- Generate \tilde{N} independent proposals $\theta^i \sim p(\theta)$, and corresponding data sets $y^\star \sim p(y|\theta^i)$ from the computer simulator $p(y|\theta)$;
- Compute the summary statistics $s^i = S(y^i)$ for each $i = 1, \dots, \tilde{N}$;
- Compute the distances $\Delta(s^i, s^{\text{obs}})$ for each $i = 1, \dots, \tilde{N}$.
- Retain proposals θ^i corresponding to those $\Delta(s^i, s^{\text{obs}})$ such that $\Delta(s^i, s^{\text{obs}}) \leq \epsilon$, for some $\epsilon \geq 0$;
- We sample from $p_{\text{ABC}}^\epsilon(\theta^\star | s^{\text{obs}})$.

ABC rejection sampling for the Beta-Binomial model

Model

$$\begin{aligned} y &\sim \text{Binomial}(m, p), \\ p &\sim \text{Beta}(\alpha, \beta) \end{aligned}$$

```

In [36]: # Define the model and the prior distribution.

Random.seed!(12) # fix random numbers

# model parameters
m = 4; n = 5; p_true = 0.7

# define the data generating function
data_generator(p) = rand(Binomial(m,p),n)

# generate data
y_obs = data_generator(p_true)

# prior
 $\alpha = 2$ ;  $\beta = 2$ 
prior = Beta( $\alpha$ , $\beta$ );

```

ABC rejection sampling for the Beta-Binomial model

```

In [38]: # ABC rejection sampling algorithm
function abc_rs(;N_proposals::Int,  $\epsilon$ ::Real)

    abc_posterior_samples = zeros(N_proposals)
    nbr_accepted_proposals = 0

    for i in 1:N_proposals

        p_star = rand(prior) # sample parameter proposal from prior
        y_star = data_generator(p_star) # generate data from the data model
         $\Delta$  = sum(abs.(sort(y_star)-sort(y_obs))) # compute ABC distance

        if  $\Delta \leq \epsilon$  # accept proposal
            nbr_accepted_proposals += 1
            abc_posterior_samples[nbr_accepted_proposals] = p_star
        end
    end

    return abc_posterior_samples[1:nbr_accepted_proposals]
end;

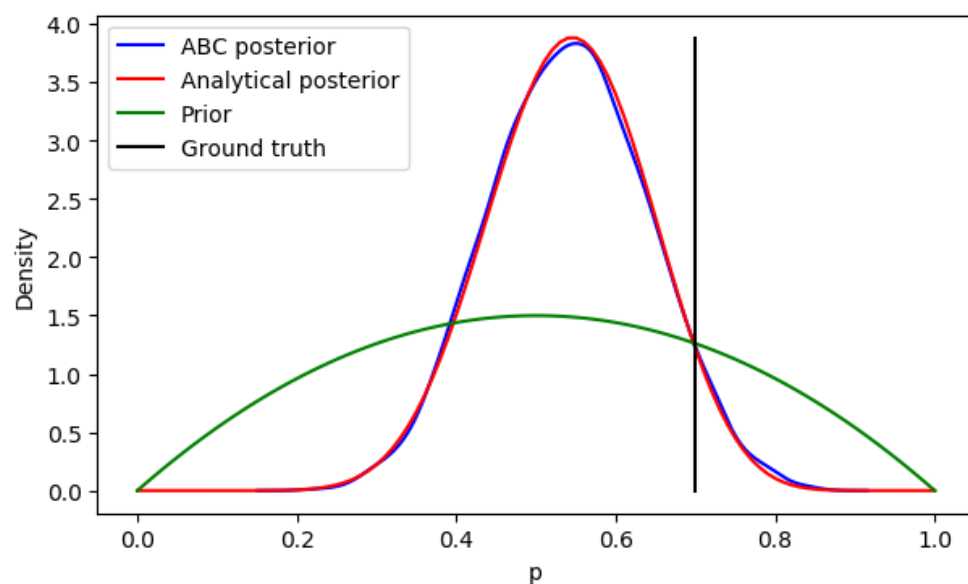
```

ABC rejection sampling for the Beta-Binomial model

```
In [56]: # Run ABC rejection sampling
abc_posterior_samples = abc_rs(N_proposals = 10^6,  $\epsilon$  = 0);
@printf "Acceptance rate: %.2f %%" length(abc_posterior_samples)/1
0^6*100
```

Acceptance rate: 0.35 %

```
In [57]: # plot posterior inference results
plot_abc_inference_results(abc_posterior_samples);
```



ABC rejection sampling for the Beta-Binomial model (with summary statistics)

```
In [43]: # define the summary statistics
S(y) = sum(y); # canonical statistic, i.e. the statistic is sufficient!
```

```

In [46]: # ABC rejection sampling algorithm
function abc_rs_summary_stats(;N_proposals::Int,  $\epsilon$ ::Real, S::Function)

    abc_posterior_samples = zeros(N_proposals)
    nbr_accepted_proposals = 0

    for i in 1:N_proposals

        p_star = rand(prior) # sample parameter proposal from prior
        y_star = data_generator(p_star) # generate data from the data model
         $\Delta$  = abs(S(y_star)-S(y_obs)) # compute ABC distance

        if  $\Delta$  <=  $\epsilon$  # accept proposal
            nbr_accepted_proposals += 1
            abc_posterior_samples[nbr_accepted_proposals] = p_star
        end
    end

    return abc_posterior_samples[1:nbr_accepted_proposals]
end;

```

ABC rejection sampling for the Beta-Binomial model (with summary statistics)

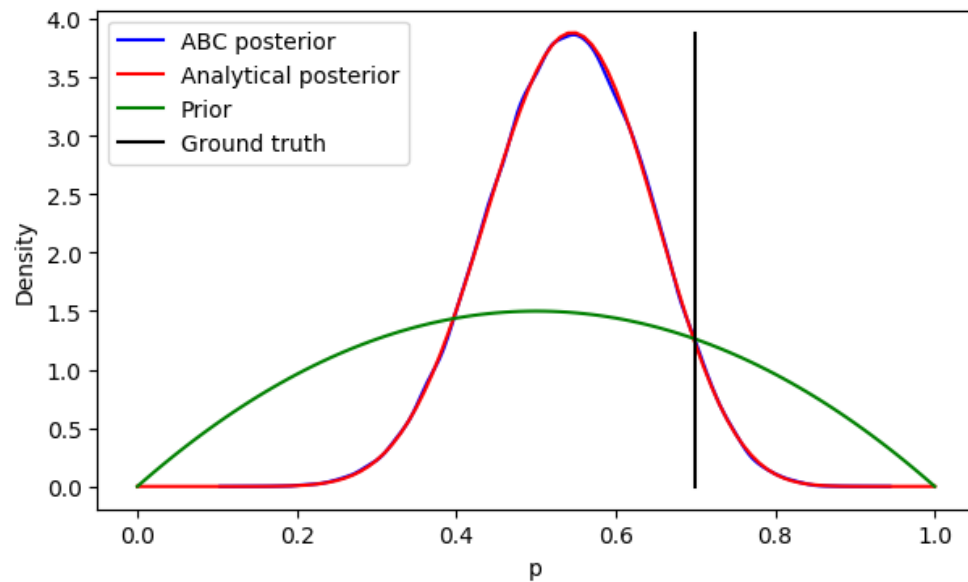
```

In [62]: # Run ABC rejection sampling
abc_posterior_samples = abc_rs_summary_stats(N_proposals = 10^6,  $\epsilon$  = 0.1, S=S);
@printf "Acceptance rate: %.2f %" length(abc_posterior_samples)/10^6*100

```

Acceptance rate: 6.77 %

```
In [63]: # plot posterior inference results
plot_abc_inference_results(abc_posterior_samples);
```



How to select/learn summary statistics

- The summary statistics *should* be low-dimensional and informative for the parameters (in the ideal case sufficient).
- The problem of selecting informative summary statistics is the main challenge when applying ABC in practice;
- Usually, summary statistics are ad-hoc and "handpicked" out of subject-domain expertise. For example for dynamic models, summaries can be autocorrelations, cross-covariances, stationary mean. For i.i.d. data could be quantiles, mean and standard deviation etc.;
- Several methods to learn/select summary statistics have been developed (see (Prangle, 2015) for a review on these methods);

Learning summary statistics using linear regression

- An important paper is Fearnhead & Prangle, 2012 where they use linear regression to learn summary statistics, they also show that the posterior mean is the *best* (in terms of loss for the posterior mean) summary statistic;
- The semi-automatic ABC (the method from Fearnhead, 2012):
 - We can sample a set of parameter-data pairs $(\theta^i, y^i)_{1 \leq i \leq N}$, by sampling θ^i from the prior, and then simulate corresponding data set y^i from the simulator $p(y|\theta)$;
 - Learn the posterior mean from the N simulations, using a linear regression model:

$$\theta_j^i = E(\theta_j|y^i) + \xi_j^i = b_{0j} + b_j h(y^i) + \xi_j^i.$$
 - After fitting the linear regression model $S_j(y^\star) = \tilde{b}_{0j} + \tilde{b}_j h(y^\star)$ is the j :th summary statistics for the proposed data set y^\star .

Semi-automatic ABC for the Beta-Binomial model: Step 1: Generate data

```
In [65]: # generate parameter-data pairs
N = 1000
parameters = rand(prior, N)
data = zeros(N,5)
for i in 1:N; data[i,:] = data_generator(parameters[i]); end
```

Semi-automatic ABC for the Beta-Binomial model: Step 2: Fit linear regression model

```
In [66]: # Fit linear regression model
        β = (data'*data)\data'*parameters # ls estimation
```

```
Out[66]: 5-element Array{Float64,1}:
         0.04676202786457585
         0.04294071054160597
         0.04782352905708373
         0.05299754219974466
         0.05234867025531753
```

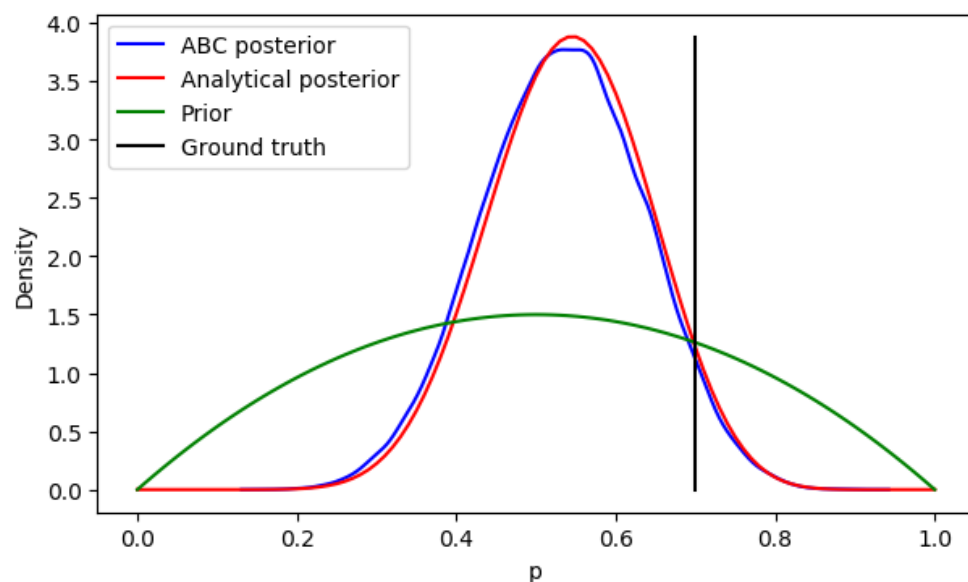
```
In [33]: # Define the new function to compute the summary statistic
        S_semi_auto(y) = y'*β;
```

Semi-automatic ABC for the Beta-Binomial model: Step 2: Run ABC algorithm

```
In [67]: # Run ABC rejection sampling
        abc_posterior_samples = abc_rs_summary_stats(N_proposals = 10^6, ε
        = 0.025, S=S_semi_auto);
        @printf "Acceptance rate: %.2f %" length(abc_posterior_samples)/1
        0^6*100
```

Acceptance rate: 7.15 %

```
In [68]: # plot posterior inference results
        plot_abc_inference_results(abc_posterior_samples);
```



Replacing linear regression with multilayer perceptron (MLP) network

- In Jiang et al., 2017 they replace the linear regression model with a MLP network, thus they have following regression model:

$$\theta^i = E(\theta|y^i) + \xi^i = f_{\beta}(y^i) + \xi^i.$$

Where f_{β} is the MLP parameterized by the weights β .

Source: (Jiang et al., 2017).

The partially exchangeable network (PEN)

- Markovian data is *partially exchangeable* (which is a property that characterizes Markovian data the same way as *exchangeability* characterizes i.i.d data);
- Now, PEN is designed such that it is invariant to the partial exchangeability property of Markovian data;
- We can write the PEN regression model as:

$$\theta^i = E(\theta|y^i) + \xi^i = \rho_{\beta_{\rho}} \left(y_{1:d}^i, \sum_{l=1}^{M-d} \phi_{\beta_{\phi}}(y_{l:l+d}^i) \right) + \xi^i.$$

- The advantage of PEN is that the architecture leverages the partial exchangeability property of Markovian data, and thus do not have to *learn* this property.

Results for the AR2 model

Model:

$$y_l = \theta_1 y_{l-1} + \theta_2 y_{l-2} + z_l, \quad z_l \sim N(0, 1).$$

Conclusions

- A practical introduction to ABC for a simple model (the Beta-Binomial model)
- We have studied how deep learning methods can be used to learn summary statistics for ABC, and presented the entire workflow for a simple example;
- PEN is particular useful to use for timeseries data since the network leverage the Markovina structure of the data.
- In Wiqvist et al. 2019 we show how it is possible to obtain good results for several models, including non-Markovian data (for many more details and examples see (Wiqvist et al. 2019)).

The end

Homepage: <http://www.maths.lu.se/staff/samuel-wiqvist/>
(<http://www.maths.lu.se/staff/samuel-wiqvist/>)

Slides: <https://github.com/SamuelWiqvist/bayesatlund2019presentation> (<https://github.com/SamuelWiqvist/bayesatlund2019presentation>)

Github: SamuelWiqvist

Twitter: samuel_wiqvist

References

Fearnhead, P. and Prangle, D. *Constructing summary statistics for approximate bayesian computation: semi-automatic approximate Bayesian computation*. Journal of the Royal Statistical Society: Series B, 74(3):419–474, 2012.

Jiang, B., Wu, T.-y., Zheng, C., and Wong, W. H. *Learning summary statistic for approximate Bayesian computation via deep neural network*. Statistica Sinica, pp. 1595–1618, 2017.

Prangle, D. *Summary statistics in approximate Bayesian computation*. arXiv:1512.05633, 2015.

Wiqvist, S., Mattei P-A., Picchini U., and Frellsen J. *Partially Exchangeable Networks and Architectures for Learning Summary Statistics in Approximate Bayesian Computation*, arXiv:1901.10230, 2019.

Extra slide: Approximate Bayesian Computation: Approximate posterior

- The joint distribution of accepted parameter-data pairs (θ^*, s^*) is

$$p(\theta^*, s^*) = p(s^* | \theta^*) p(\theta^*) I(\Delta(s^*, s^{\text{obs}}) \leq \epsilon),$$

where $s^{\text{obs}} = S(y^{\text{obs}})$, I indicator kernel, Δ distance function, and ϵ the threshold.

- Now assume that $S(s^*) = S(s^{\text{obs}})$ iff $y^* = y^{\text{obs}}$ and let $\epsilon = 0$. Now marginalizing s^* yields the true posterior:

$$p(\theta^*) = \int p(s^* | \theta^*) p(\theta^*) I(\Delta(s^*, s^{\text{obs}}) \leq \epsilon) ds^* = p(y | \theta^*) p(\theta^*)$$

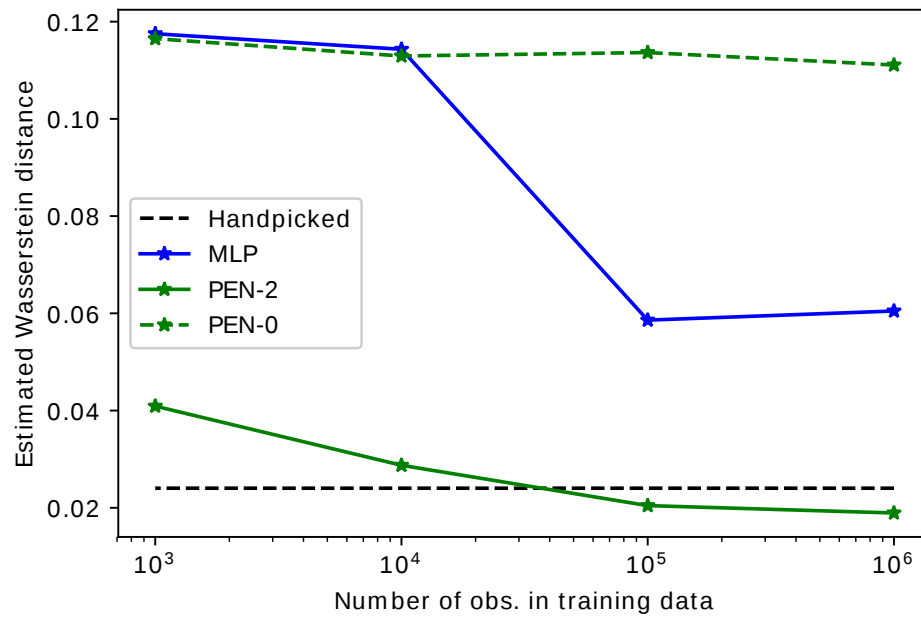
- However, in (almost) all situations we sample from an approximate posterior:

$$p_{\text{ABC}}^{\epsilon}(\theta^* | s^{\text{obs}}) \propto \int p(s^* | \theta^*) p(\theta^*) I(\Delta(s^*, s^{\text{obs}}) \leq \epsilon) ds^*.$$

Extra slide: Results for the AR2 model with observation noise

Model:

$$y_l = \theta_1 y_{l-1} + \theta_2 y_{l-2} + z_l, \quad z_l \sim N(0, 1).$$

**Extra slide: Results for the MA2 model with observation noise**

Model:

$$y_i = x_i + e_i, \quad e_i \sim N(0, \sigma_e),$$

$$x_i = z_i + \theta_1 z_{i-1} + \theta_2 z_{i-2}, \quad z_i \sim N(0, 1).$$

